

Introduction

Assignment #1 is worth 10% of your final mark.

Assignment #2 will be worth 15% of your final mark.

This assignment is new and has never been used in CMPUT 274 before. Therefore, on the one hand, there may be a higher rate of corrections, clarifications, and additional information provided than the mature Weekly Exercises that you have seen so far. On the other hand, the Python code that you wrote for some of the Weekly Exercises can (by design) be re-used for this assignment.

Also, an assignment of this scope includes many unspecified details for which a **reasonable design decision** can be made (i.e., does **not** contradict an existing specification), stated explicitly (e.g., in a comment in the code), and implemented.

This assignment will require you to know something about:

1. Machine-learned classifiers
2. Python programming
3. Classes and object-oriented programming in Python
4. Preprocessing and stopwords
5. Word frequency
6. Cross validation
7. Confusion matrices

For this assignment, you will be extending the object-oriented (OO) classifier (`ooclassifierbase.py`) program provided to you. Your modified program will be renamed `ooclassifier.py` (see Submission Guidelines). You will be adding preprocessor functionality (similar to Weekly Exercise #4), adding a new classification algorithm (based on frequent words, similar to Weekly Exercise #3), and adding the ability to support n-fold cross-validation.

The main challenge of this assignment is **not** in the total number of lines of code that has to be implemented, but rather in understanding the original code (`ooclassifierbase.py`) well enough to be able to extend it with new functionality. To be sure, there is a non-trivial amount of new Python code that must be written, but this assignment is not only about writing code.

Task I: Preprocessing in the Classifier

1. Add a new method `preprocess_words(mode='')` to class `TrainingInstance`.

When invoked, `preprocess_words(mode='')` applies all of the preprocessing actions from Weekly Exercise #4 Text Preprocessor to all the words in that particular training instance object.

2. Add a new method `preprocess(mode='')` to class `TrainingSet`.

When invoked, `preprocess(mode='')` will use `preprocess_words(mode='')` in class `TrainingInstance` to perform preprocessing for all training instances in a particular training dataset.

3. You may add additional methods and attributes to class `TrainingSet` and class `TrainingInstance`, as long as `preprocess(mode='')` and `preprocess_words(mode='')` work as required.
4. Do not change the interface or semantics of any of the existing methods and attributes of the existing classes.

When `preprocess(mode='')` returns, **all** of the training instances in the **entire** training dataset will have been preprocessed according to the specifications in Weekly Exercise #4 Text Preprocessor. The full description of the preprocessing steps are as described in section “Your Task #1: Full Preprocessing” of Weekly Exercise #4.

HINT (Oct. 23, 2020): *After preprocessing, the training instances still need to be explicitly re-classified. In computing science, preprocessing and processing (e.g., classifying) are usually separate steps, because there may be multiple steps or kinds of “preprocessing” (e.g., doing some but not all modes of preprocessing; see below) before the data is actually used.*

Both new methods take a default argument called `mode` that is a string. If the string is empty, then it performs the full preprocessing from “Your Task #1: Full Preprocessing” of Weekly Exercise #4.

When the string `mode` is one of the values `'keep-digits'`, `'keep-stops'` or `'keep-symbols'`, the methods behave according to the description in section “Your Task #2: Optional Command Line Argument” of Weekly Exercise #4. Whereas in Weekly Exercise #4 the mode came from the command line, in this assignment the mode parameter comes from the call-site(s) of `preprocess_words()` and `preprocess()`.

Task II: Word Frequency Classification Algorithm

So far, the classification algorithm based on hard-coded target words works well enough, but there is actually no machine learning in the algorithm. The target words are hard-coded beforehand and do not change based on the input or training dataset, which is necessary for machine learning.

Therefore:

1. Create a new class `ClassifyByTopN`, which is a subclass of class `ClassifyByTarget`.
2. The class `ClassifyByTopN` includes a new method `target_top_n(tset, num=5, label='')` which replaces the current list of target words with a new list of target words based on word frequency. Details below.
3. You may add additional methods and attributes to class `ClassifyByTopN`, as long as `target_top_n(tset, num=5, label='')` works as required.
4. Do not change the interface or semantics of any of the existing methods and attributes of the existing classes.

Specifically, `target_top_n(tset, num=5, label='')` counts all of the words, in all of the training instances, whose label match the string `label`, of object `tset` which is of class `TrainingSet`.

Of course, if `tset` has been preprocessed, then the word frequency count will be based on the words remaining after the preprocessing is done. Otherwise, the word frequency count is based on whatever words are currently in the training dataset.

After the word frequency count, the top `num` most frequent words become the new target words list. The default value for `num` is 5. If there is a tie for the count at the `num`-th rank, then **all** of the words with the same count are also included in the new list of target words. Therefore, it is possible for the number of target words to be larger than `num`, due to ties in the counts.

Since class `ClassifyByTopN`, is a subclass of class `ClassifyByTarget`, it should support all of the existing methods of class `ClassifyByTarget`, such as (in particular, but not a complete list) `classify()`, `eval_training_set()`, and `print_confusion_matrix()`.

Task III: Create N-folds for Cross Validation

N-fold cross validation is an important technique in evaluating machine learning (ML) algorithms and training data, such as classifiers. If an ML model trains on data, and then that same data is used to evaluate the model, one could simply be evaluating how well the model memorizes (see memoization) the answers it has seen before.

By analogy, it is like giving out all the questions and answers to the final exam beforehand, and then only using those questions on the actual final exam. Such an exam would be testing your memory, not your understanding.

In n -fold cross validation, a portion of the training data is withheld, the rest of the training data is used, then the withheld portion is used to test classifier. Specifically, the whole training dataset is divided into n different partitions or folds. When combined (e.g., set union), the n folds recreate the original training dataset (possibly with some changes in ordering).

1. Create a new method `return_nfolds(num=3)` in class `TrainingSet`, that returns a list of `num` objects of class `TrainingSet`. Each of the objects returned contains a partition or **fold** of the original training dataset. Each of the objects should contain a deepcopy of all attributes of the `TrainingSet`. The default value for `num` is 3.
2. Create a new method `copy()` in class `TrainingSet` that returns (after making a deepcopy) an object of class `TrainingSet` that contains the same attributes (e.g., training instances) as the original object of class `TrainingSet`. This `copy()` is similar to `mylist.copy()`, except that `deepcopy` (instead of a shallow copy) is used.
3. Create a new method `add_fold(tset)` in class `TrainingSet` that adds (via `deepcopy`) the training instances of `tset` (which is of class `TrainingSet`) to an object of class `TrainingSet`.
4. You may add additional methods and attributes to class `TrainingSet`, as long as methods `return_nfolds(num=3)`, `copy()`, and `add_fold(tset)` work as required.
5. Do not change the interface or semantics of any of the existing methods and attributes of the existing classes.

Do not be concerned with any time or space efficiency issues related to the use of `deepcopy`. Those concerns can be addressed by future optimizations.

A basic **round robin** (or interleaved) strategy should be used to create the folds. It is possible that the folds are not exactly equal in size, as they can be off-by-one depending on whether num divides evenly into the number of training instances.

In combination with class `ClassifyByTopN`, proper n-fold cross validation experiments can now be performed.

Sample Output:

In a few days, sample output will be provided on eClass.

Submission Guidelines:

Submit all of the required files (and no other files) as **one** properly formed compressed archive called either `ooclassifier.tar.gz`, or `ooclassifier.tgz`, or `ooclassifier.zip` (for full marks, please do **not** use `.rar`):

- when your archive is extracted, it should result in exactly *one directory* called `ooclassifier` (use this exact name) with the following files in that directory:
- `ooclassifier.py` (use this exact name) contains all of your Python code and *docstrings-based documentation*. NOTE: We will be testing your code via `from ooclassifier import *` or similar.
- your `README` (use this exact name) conforms with the Code Submission Guidelines.
- No other files should be submitted.

Note that your files and functions must be named **exactly** as specified above.

A new tool has been developed by the TAs to help check and validate the format of your `tar` or `zip` file *prior* to submission. To run it, you will need to download it into the VM, and place it in the same directory as your compressed archive (e.g., `ooclassifier.zip`).

You can read detailed instructions and more explanation about this new tool in Submission Validator: Instructions (at the top of the Weekly Exercises tab), or run:

```
python3 submission_validator.py --help
```

after you have downloaded the script to see abbreviated instructions printed to the terminal.

If your submission passes this validation process, and all validation instructions have been followed properly, you will not lose any marks related to the format of your submission. (Of course, marks can still be deducted for correctness, design, and style reasons, but not for submission correctness.)

When your marked assignment is returned to you, there is a 7-day window to request the reconsideration of any aspect of the mark. After the window, we will only change a mark if there is a clear mistake on our part (e.g., incorrect arithmetic, incorrect recording of the mark). At any time during the term, you can request additional feedback on your submission.